

Search-based Path Finding

Hongkai YE

2021.10.05



Outline



1. Graph Search Basis



2. Dijkstra and A*



3. Jump Point Search

Graph Search Basis



Configuration Space




Configuration Space

- **Robot configuration:** a specification of the positions of all points of the robot
- **Robot degree of freedom (DOF):** The minimum number n of real-valued coordinates needed to represent the robot configuration
- **Robot configuration space:** a n -dim space containing all possible robot configurations, denoted as **C-space**
- Each robot pose is a **point** in the **C-space**



Configuration Space Obstacle

- Planning in workspace
 - Robot has different shape and size
 - Collision detection requires knowing the robot geometry - time consuming and hard



(1) Rectangular mobile robot

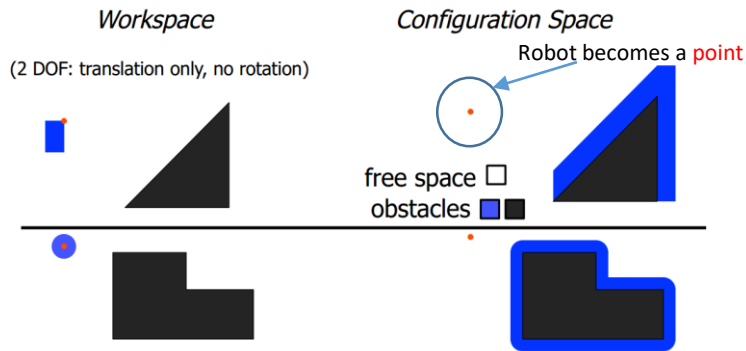


(2) Circular mobile robot



Configuration Space Obstacle


- Planning in configuration space
 - Robot is represented by a point in C-space, e.g. position (a **point** in R^3), pose (a **point** in $SO(3)$), etc.
 - Obstacles need to be represented in configuration space (one-time work prior to motion planning), called configuration space obstacle, or C-obstacle
 - C-space = (C-obstacle) \cup (C-free)
 - The path planning is finding a path between start **point** q_{start} and goal **point** q_{goal} within C-free





Workspace and Configuration Space Obstacle

- In workspace
 - Robot has shape and size (i.e. hard for motion planning)
- In configuration space: C-space
 - Robot is a point (i.e. easy for motion planning)
 - Obstacle are represented in C-space prior to motion planning
- Representing an obstacle in C-space can be extremely complicated. So approximated (but more conservative) representations are used in practice.



If we model the robot conservatively as a ball with radius δ_r , then the C-space can be constructed by inflating obstacle at all directions by δ_r .


Graph and Search Method




Search-based Method

Graphs


Graphs have **nodes** and **edges**




Undirected



Weighted



Directed

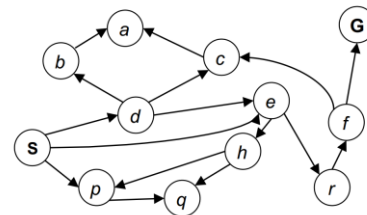




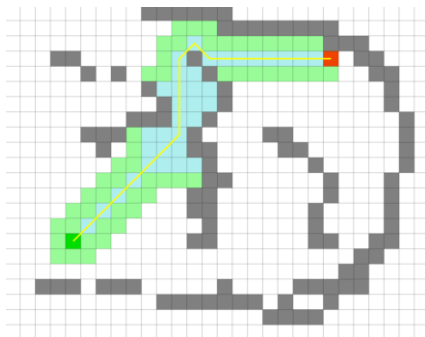
Search-based Method

- State space graph: a mathematical representation of a **search algorithm**

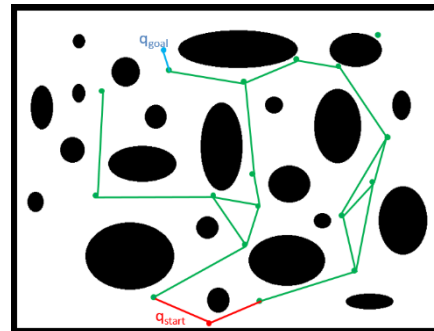
- For every search problem, there's a corresponding state space graph
- Connectivity between nodes in the graph is represented by (directed or undirected) edges



*Ridiculously tiny search graph
for a tiny search problem*



Grid-based graph: use grid as vertices and grid connections as edges




The graph generated by probabilistic roadmap (PRM)



Graph Search Overview

- The search always start from start state X_S
 - Searching the graph produces a search tree
 - Back-tracing a node in the search tree gives us a path from the start state to that node
 - For many problems we can never actually build the whole tree, too large or inefficient – we only want to reach the goal node asap.





Graph Search Overview

- Maintain a **container** to store all the nodes **to be visited**
- The container is initialized with the start state X_s
- Loop
 - **Remove** a node from the container according to some pre-defined score function
 - Visit a node
 - **Expansion**: Obtain all **neighbors** of the node
 - Discover all its neighbors
 - **Push** them (**neighbors**) into the container
- End Loop



Graph Search Overview


- Question 1: When to end the loop?
 - Possible option: End the loop when the container is empty
- Question 2: What if the graph is cyclic?
 - When a node is removed from the container (expanded / visited), it should never be added back to the container again
- Question 3: In what way to remove the right node such that the **goal state can be reached as soon as possible**, which results in less expansion of the graph node.



Graph Traversal


- Breadth First Search (BFS) vs. Depth First Search (DFS)

BFS uses "first in first out"



This is a **queue**

DFS uses "last in first out"





This is a **stack**



Depth First Search (DFS)

- Strategy: remove / expand the deepest node in the container





-





Depth First Search (DFS)


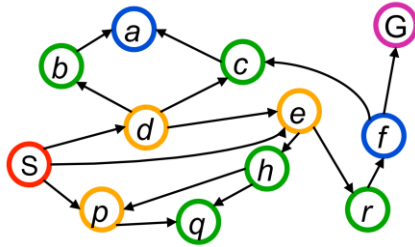


Courtesy: Amit Patel's Introduction to A*, Stanford



Breadth First Search (BFS)

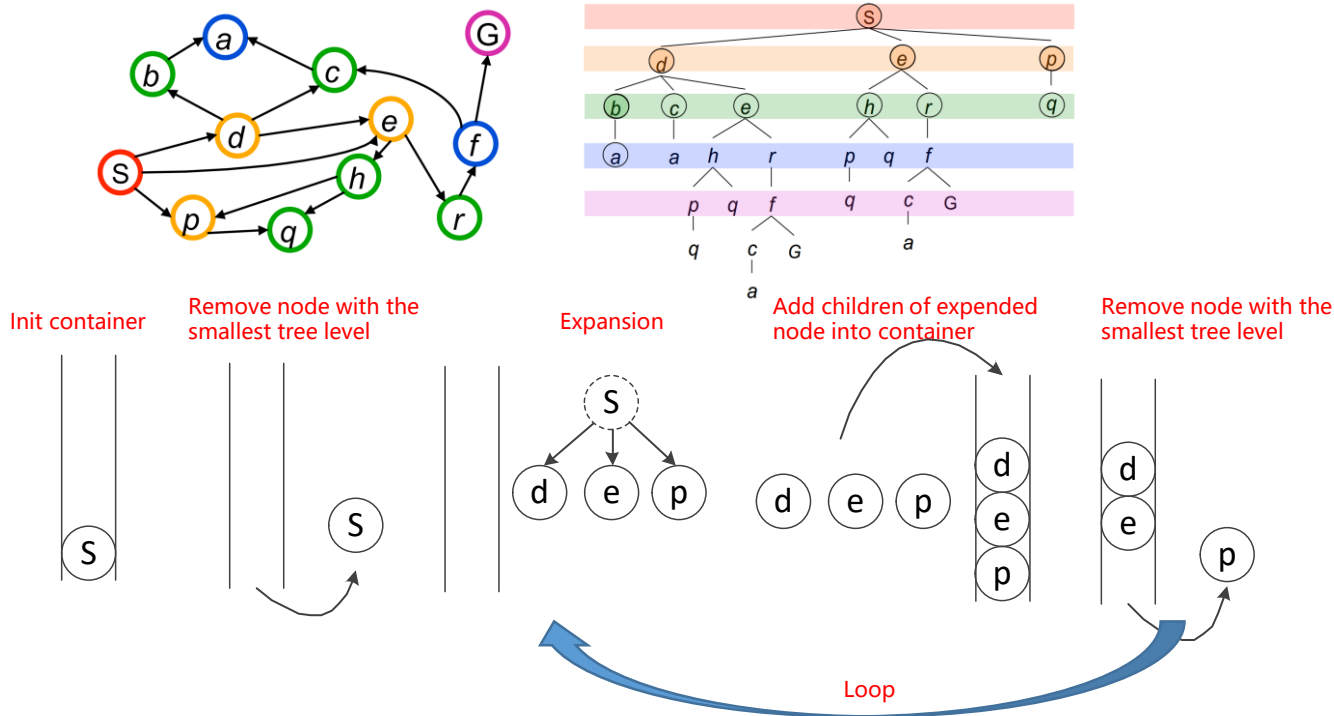
- Strategy: remove / expand the shallowest node in the container





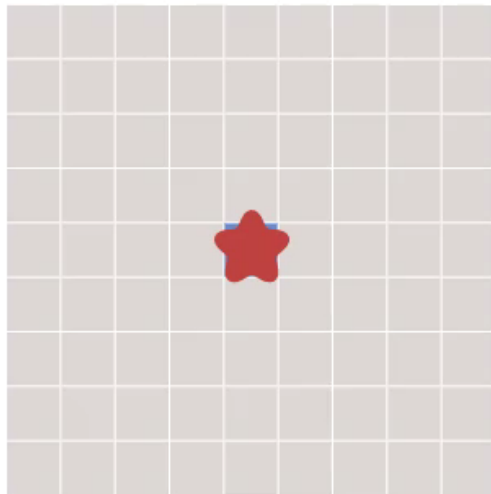
Breadth First Search (BFS)

- Implementation: maintain a first in first out (FIFO) container (i.e. queue)





Breadth First Search (BFS)

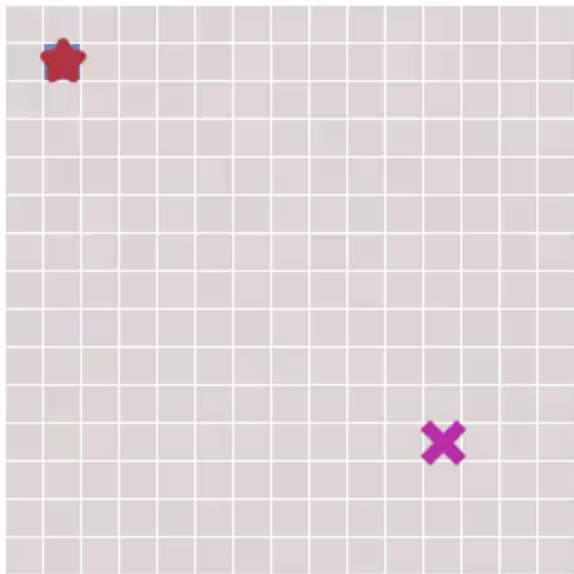


Courtesy: Amit Patel's Introduction to A*, Stanford

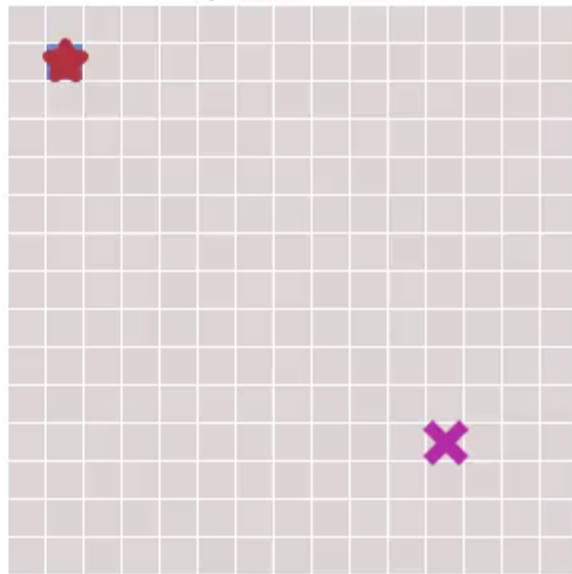


BFS vs. DFS: which one is useful?

Breadth First Search



Depth First Search



Remember BFS.




Heuristic search



Greedy Best First Search

- BFS and DFS pick the next node off the frontiers based on which was "first in" or "last in".
- Greedy Best First picks the "best" node according to some rule, called a **heuristic**.
- **Definition:** A heuristic is a **guess** of how close you are to the target.

- A heuristic guides you in the right direction.
- A heuristic should be easy to compute.



- **Euclidean Distance**
- **Manhattan Distance**

Both are approximations for the actual **shortest path**.



Greedy Best First Search



Looks pretty good.



Greedy Best First Search

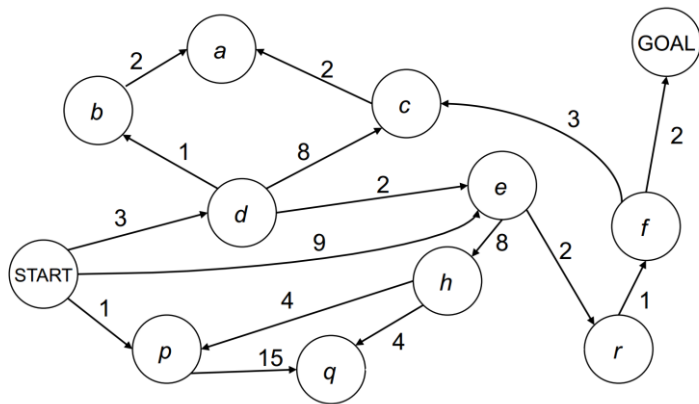


But with obstacles ...



Costs on Actions

- A practical search problem has a **cost “C”** from a node to its neighbor
 - Length, time, energy, etc.
- When all weight are 1, BFS finds the optimal solution
- For general cases, how to find the **least-cost path** as soon as possible?



Dijkstra and A*

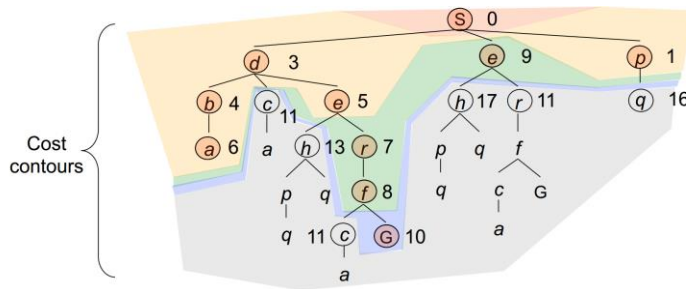
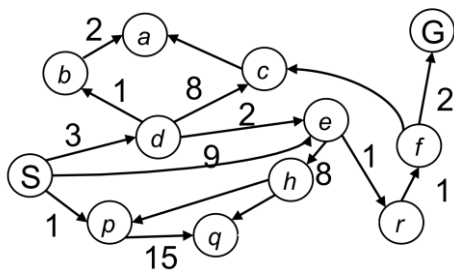


Algorithm Workflow



Dijkstra's Algorithm

- Strategy: expand/visit the node with **cheapest accumulated cost $g(n)$**
 - $g(n)$: The current best estimates of the accumulated cost from the start state to node "n"
 - Update the accumulated costs $g(m)$ for all unexpanded neighbors "m" of node "n"
 - A node that has been expanded/visited is guaranteed to have the smallest cost from the start state




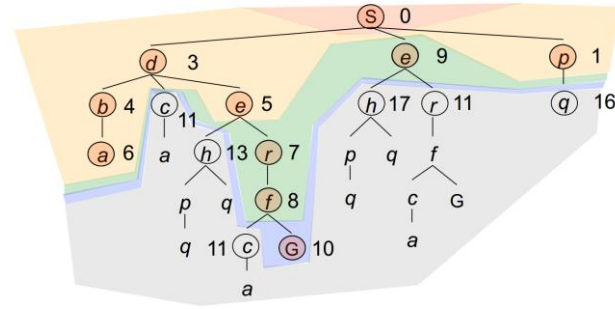
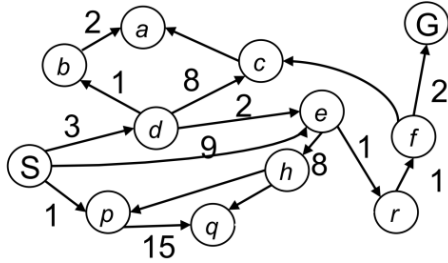


Dijkstra's Algorithm

- Maintain a **priority queue** to store all the nodes to be expanded
- The priority queue is initialized with the start state X_s
- Assign $g(X_s)=0$, and $g(n)=\text{infinite}$ for all other nodes in the graph
- Loop
 - If the queue is empty, return FALSE; break;
 - Remove the node "n" with the lowest $g(n)$ from the priority queue
 - Mark node "n" as expanded
 - If the node "n" is the goal state, return TRUE; break;
 - For all unexpanded neighbors "m" of node "n"
 - If $g(m) = \text{infinite}$
 - $g(m) = g(n) + C_{nm}$
 - Push node "m" into the queue
 - If $g(m) > g(n) + C_{nm}$
 - $g(m) = g(n) + C_{nm}$
 - end
- End Loop



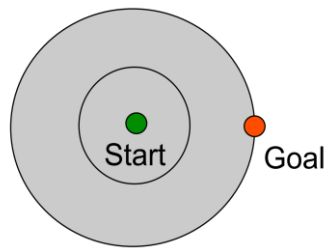
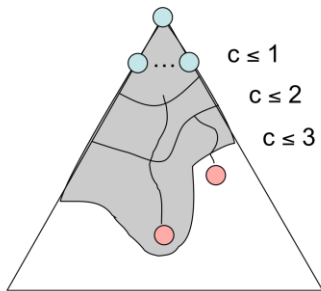
Dijkstra's Algorithm





Pros and Cons of Dijkstra's Algorithm

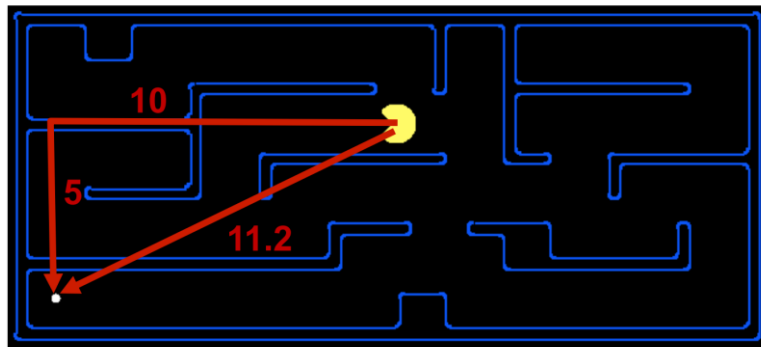
- The good:
 - Complete and optimal
- The bad:
 - Can only see the cost accumulated so far (i.e. the uniform cost), thus exploring next state in every "direction"
 - No information about goal location





Search Heuristics

- Recall the heuristic introduced in **Greedy Best First Search**
- Overcome the shortcomings of uniform cost search by **inferring the least cost to goal (i.e. goal cost)**
- Designed for particular search problem
- Examples: Manhattan distance VS. Euclidean distance





A*: Dijkstra with a Heuristic

- Accumulated cost
 - $g(n)$: The current best estimates of the accumulated cost from the start state to node “n”
- Heuristic
 - $h(n)$: The **estimated least cost** from node n to goal state (i.e. goal cost)
- The least estimated cost from start state to goal state passing through node “n” is $f(n) = g(n) + h(n)$
- Strategy: expand the node with **cheapest $f(n) = g(n) + h(n)$**
 - Update the accumulated costs $g(m)$ for all unexpanded neighbors “m” of node “n”
 - A node that has been expanded is guaranteed to have the smallest cost from the start state




A* Algorithm

- Maintain a **priority queue** to store all the nodes to be expanded
- The heuristic function $h(n)$ for all nodes are pre-defined
- The priority queue is initialized with the start state X_s
- Assign $g(X_s)=0$, and $g(n)=\text{infinite}$ for all other nodes in the graph
- Loop
 - If the queue is empty, return FALSE; break;
 - **Remove** the node "n" with the lowest $f(n)=g(n)+h(n)$ from the priority queue
 - Mark node "n" as **expanded**
 - If the node "n" is the goal state, return TRUE; break;
 - For all **unexpanded** neighbors "m" of node "n"
 - If $g(m) = \text{infinite}$
 - $g(m) = g(n) + C_{nm}$
 - Push node "m" into the queue
 - If $g(m) > g(n) + C_{nm}$
 - $g(m) = g(n) + C_{nm}$
 - end
- End Loop


Only difference comparing to
Dijkstra's algorithm




A* Example




Initial queue




1st dequeue
& expand




1st enqueue




2nd dequeue
& expand



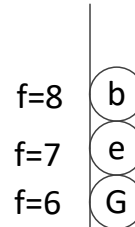
2nd enqueue




3rd dequeue
& expand



3rd enqueue




4th dequeue





A* Optimality




- What went wrong?
- For node A: actual least cost to goal (i.e. goal cost) < estimated least cost to goal (i.e. heuristic)
- We need the estimate to be **less than** actual least cost to goal (i.e. goal cost) **for all nodes**!




Admissible Heuristics

- A Heuristic h is **admissible** (optimistic) if:
 - $h(n) \leq h^*(n)$ for every node " n ", where $h^*(n)$ is the true least cost to goal from node " n "
- If the heuristic is admissible, the A* search is optimal
- Coming up with admissible heuristics is most of what's involved in using A* in practice.
- Example:





A* Optimal Efficiency



State	H
S	7
A	6
B	2
C	1
G	0

- A Heuristic h is **consistent** (monotone) if:
 - $h(x) \leq d(x, y) + h(y)$ for every edge (x, y)
- Triangle inequality.
- Once a node is expanded, the cost by which it was reached is the lowest possible. If not, a node will need repeated expansion every time a new best (so-far) cost is achieved for it.
- All consistent heuristics are admissible.



Heuristic Design

An admissible heuristic function has to be designed case by case.


- Euclidean Distance
- Manhattan Distance

Is Euclidean distance (L2 norm) admissible?

Always

Is Manhattan distance (L1 norm) admissible?

Depends



Is L_∞ norm distance admissible?

Always

Is 0 distance admissible?

Always



Sub-optimal Solution

What if we intend to use an over-estimate heuristic?

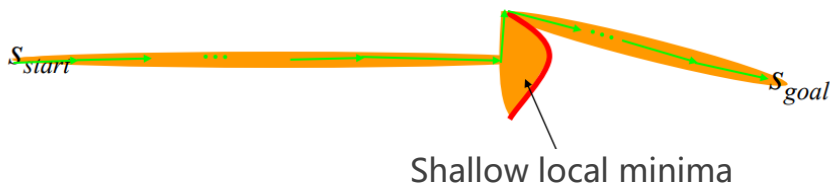


- Suboptimal path
- Faster



Weighted A*:

Expands states based on $f = g + \epsilon h, \epsilon > 1$ = bias towards states that are closer to goal.



- Weighted A* Search:

- Optimality vs. speed
- ϵ -suboptimal:
 $\text{cost}(\text{solution}) \leq \epsilon \text{cost}(\text{optimal solution})$
- It can be orders of magnitude faster than A*


Weighted A* -> Anytime A* -> ARA* -> D*

Beyond the scope of this course




Greedy Best First Search vs. Weighted A* vs. A*

$$f = a \cdot g + b \cdot h$$




Most Greedy

$$a = 0, b = 1$$



Tunable Greediness

$$a = 1, b = \varepsilon > 1$$



Optimal

$$a = 1, b = 1$$







Engineering Considerations







Example: Grid-based Path Search

How to represent grids as graphs?

Each cell is a node. Edges connect adjacent cells.



Common Choice!



4 connection

8 connection



The Best Heuristic

- Recall:

- Is Euclidean distance (L2 norm) admissible?
- Is Manhattan distance (L1 norm) admissible?
- Is L_∞ norm distance admissible?
- Is 0 distance admissible?




They are useful, but none of them is the best choice, why?

Because none of them is **tight**.

Tight means how close they measure the true shortest distance.

Euclidean Heuristic



Why so many nodes expanded?

Because Euclidean distance is far from the **truly theoretical optimal solution**.



The Best Heuristic

How to get the **truly theoretical optimal solution?**

Fortunately, the grid map is highly structural.

- You don't need to search the path.
- It has the **closed-form solution!**

```
dx=abs(node.x -goal.x)
dy=abs(node.y -goal.y)
```

```
h=D * (dx + dy) + (D2 - 2 * D) * min(dx, dy)
```

or

```
h = D * max(dx, dy) + (D2-D) * min(dx, dy)
```

or


```
if (dx > dy)
```

```
    h = D * (dx-dy) + D2 * dy
```


```
else
```

```
    h = D * (dy-dx) + D2 * dx
```

Compare



Diagonal Heuristic



Euclidean Heuristic

When $D = 1$ and $D2 = 1$, this is called the *Chebyshev distance*.

When $D = 1$ and $D2 = \sqrt{2}$, this is called the *Octile distance*

For 3D case, we also have a similar version of this.





Tie Breaker

- Many paths have the same f value.
- No differences among them making them explored by A* equally.

- Manipulate the f value breaks the tie.
- Make same f values differ.
- Interfere h slightly.

$$h = h \times (1.0 + p)$$

$$p < \frac{\text{minimum cost of one step}}{\text{expected maximum path cost}}$$



Slightly breaks the
admissibility of h , does
it matter?



Tie Breaker


Core idea of tie breaker:

Find a preference among same cost paths

- When nodes having same f , compare their h .
- Add deterministic random numbers to the heuristic or edge costs (A hash of the coordinates).
- Prefer paths that are along the straight line from the starting point to the goal.

$$\begin{aligned}dx1 &= \text{abs}(\text{node}.x - \text{goal}.x) \\dy1 &= \text{abs}(\text{node}.y - \text{goal}.y) \\dx2 &= \text{abs}(\text{start}.x - \text{goal}.x) \\dy2 &= \text{abs}(\text{start}.y - \text{goal}.y) \\cross &= \text{abs}(dx1 \times dy2 - dx2 \times dy1) \\h &= h + cross \times 0.001\end{aligned}$$

- ... Many customized ways

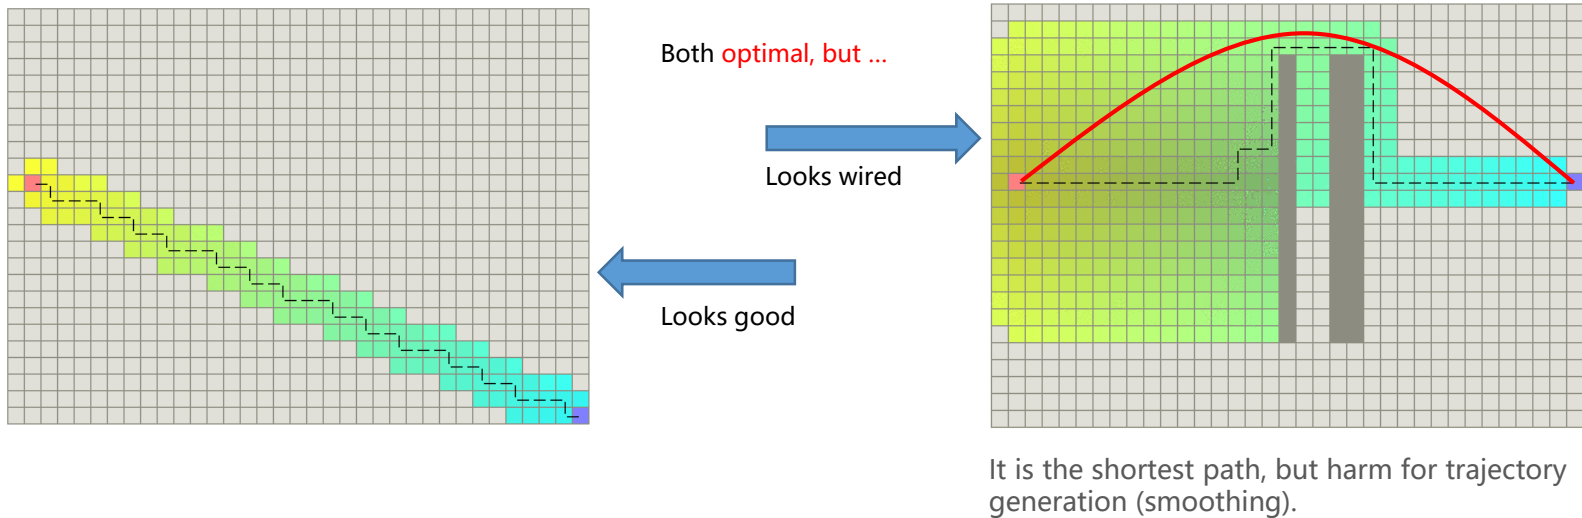


Better/other **tie breaker**?



Tie Breaker

- Prefer paths that are along the straight line from the starting point to the goal.



Or a systematic approach: **Jump Point Search (JPS)**

Jump Point Search

A method to deal with path symmetry for grid-based path search




Jump Point Search


Core idea of JPS:

Find symmetry and break them.

A* explore all symmetric path.



JPS choose one path.




Grey node: Added in the Open List.

JPS explores intelligently, because it always looks ahead based on a rule.




Look Ahead Rule

straight





diagonal



Consider:

- current node x
- x 's expanded direction



Neighbor Pruning

- Gray nodes: **inferior neighbors**, when going to them, the path without x is cheaper. Discard.
- White nodes: **natural neighbors**.
- We only need to consider natural neighbors when expand the search.


Forced Neighbors

- There is obstacle adjacent to x
- Red nodes are **forced neighbors**.
- A cheaper path from x 's parent to them is blocked by obstacle.

See: <http://users.cecs.anu.edu.au/~dharabor/data/papers/harabor-grastien-aaai11.pdf> Equation 1/2





Jumping Rules



Jumping Straight

Look Ahead Rule




Jumping Diagonally

- Recursively apply **straight** pruning rule and identify y as a **jump point successor** of x. This node is interesting because it has a neighbor z that cannot be reached optimally except by a path that visits x then y.
- Recursively apply the **diagonal pruning** rule and identify y as a **jump point successor** of x.
- Before each diagonal step we first recurse straight. Only if both straight recursions fail to identify a jump point do we step diagonally again.
- Node w, a forced neighbor of x, is expanded as normal. (also push into the open list, the **priority queue**)




Jump Point Search



- Expand horizontally and vertically.
- Both jumps end in obstacles.
- Move diagonally.




Jump Point Search



- Expand horizontally and vertically.
- Both jumps end in obstacles.
- Move diagonally.




Jump Point Search



- Expand horizontally and vertically.
- Both expansions end in obstacles.
- Move diagonally.




Jump Point Search



- Vertically expansion end in obstacle.
- Right-ward expansion finds a node with a **forced neighbor**.

Recall the rule

1		3
4	x	5
6	7	8




So we have

1	2	3
4	x	5
6		8



Jump Point Search



- Now this node is of interested.
- Put it to open list.




Jump Point Search


Recall A*'s pseudo-code, JPS's is all the same!

- Maintain a priority queue to store all the nodes to be expanded
- The heuristic function $h(n)$ for all nodes are pre-defined
- The priority queue is initialized with the start state X_s
- Assign $g(X_s)=0$, and $g(n)=\text{infinite}$ for all other nodes in the graph
- Loop
 - If the queue is empty, return FALSE; break;
 - Remove the node "n" with the lowest $f(n)=g(n)+h(n)$ from the priority queue
 - Mark node "n" as expanded
 - If the node "n" is the goal state, return TRUE; break;
 - For all unexpanded neighbors "m" of node "n"
 - If $g(m) = \text{infinite}$
 - $g(m) = g(n) + C_{nm}$
 - Push node "m" into the queue
 - If $g(m) > g(n) + C_{nm}$
 - $g(m) = g(n) + C_{nm}$
 - end
- End Loop

A*: "Geometric" neighbors



JPS: "Jumping" neighbors






Example




Jumping Example

Planning Case






Jumping Example



- Expand—> move diagonally
- Find a critical node finally, add it into open list.
- Pop it (the only one) from the open list.
- Expand vertically, end at obstacles.




Jumping Example



- Expand horizontally, meets a node with a forced neighbor.
- Add it to open list

Recall the rule

1		3
4	x	5
6	7	8




So we have

1	2	3
4	x	5
6		8




Jumping Example



- Expand diagonally, expand, find nothing.
- Finish the expansion of the current node.



Jumping Example




- Examine the “new best” node in the open list.
- Expand horizontally.
- Finds nothing.

Remember the rule

1	2	3
4 → x		5
6		8




Jumping Example



- Move diagonally.
- Expand along vertical and horizontal first.




Jumping Example



- Finds nothing.
- Move diagonally.




Jumping Example



- Expand horizontally and vertically.
- Finds the goal. Equally interested as finding a node with a forced neighbor.
- Add this node to open list.
- Finish the expand of the current node (No naturally neighbors left).
- Pop it out of the open list.



Jumping Example




- Examine the “new best” node in the open list.
- Expand horizontally (nowhere), and vertically (finds the goal).
- The end.




Jumping Example

Final Path






Example



(1)



(2)


Thanks:

<https://zerowidth.com/2013/a-visual-explanation-of-jump-point-search.html>




Is JPS always better?

Maze-like environments



JPS




A*

Do more tests by yourself!

Thanks: <http://qiao.github.io/PathFinding.js/visual/>



Is JPS always better?



- This is a simple example saying “No.”
- This case may commonly occur in robot navigation.
- Robot with limited FOV, but a global map/large local map.

Conclusion:

- Most time, especially in complex environments, JPS is better, but far away from “always”. Why?
- JPS reduces the number of nodes in **Open List**, but increases the number of **status query**.
- **JPS's limitation**: only applicable to uniform grid map.



Thanks for Listening

